

Homework 4

*Instructor: Vatsal Sharan**Due: November 16 by 2:00 pm PST*

A reminder on collaboration policy and academic integrity: Our goal is to maintain an optimal learning environment. You can discuss the homework problems at a high level with other groups, but you should not look at any other group's solutions. Trying to find solutions online or from any other sources for any homework or project is prohibited, will result in zero grade and will be reported. To prevent any future plagiarism, uploading any material from the course (your solutions, quizzes etc.) on the internet is prohibited, and any violations will also be reported. Please be considerate, and help us help everyone get the best out of this course.

Please remember the Student Conduct Code (Section 11.00 of the USC Student Guidebook). General principles of academic honesty include the concept of respect for the intellectual property of others, the expectation that individual work will be submitted unless otherwise allowed by an instructor, and the obligations both to protect one's own academic work from misuse by others as well as to avoid using another's work as one's own. All students are expected to understand and abide by these principles. Students will be referred to the Office of Student Judicial Affairs and Community Standards for further review, should there be any suspicion of academic dishonesty.

Notes on notation:

- Unless stated otherwise, scalars are denoted by small letter in normal font, vectors are denoted by small letters in bold font and matrices are denoted by capital letters in bold font.
- $\|\cdot\|$ means L2-norm unless specified otherwise, *i.e.*, $\|\cdot\| = \|\cdot\|_2$.

Instructions

We recommend that you use LaTeX to write up your homework solution. However, you can also scan handwritten notes. The homework will need to be submitted on Gradescope.

Theory-based Questions

Problem 1: Decision Trees (12pts)

Consider a binary dataset with 400 examples, where half of them belong to class A and the rest belong to class B. Next, consider two decision stumps (i.e. trees with depth 1) \mathcal{T}_1 and \mathcal{T}_2 , each with two children. For \mathcal{T}_1 , the left child has 150 examples in class A and 50 examples in class B. For \mathcal{T}_2 , the left child has 0 examples in class A and 100 examples in class B. (You can infer the number of examples in the right child using the total number of examples.)

1.1 (6 pts) In class, we discussed entropy and Gini impurity as measures of uncertainty at a leaf. Another possible metric is the classification error at the leaf, assuming that the prediction at the leaf is the majority class among all examples that belong to that leaf. For each leaf of \mathcal{T}_1 and \mathcal{T}_2 , compute the entropy (base e), Gini impurity and classification error. You can either exactly express the final numbers in terms of fractions and logarithms, or round them to two decimal places.

Classification error:

$$\epsilon_{1,L} = \frac{50}{150 + 50} = 0.25 \quad (0.5 \text{ point})$$

$$\epsilon_{1,R} = \frac{50}{50 + 150} = 0.25 \quad (0.5 \text{ point})$$

$$\epsilon_{2,L} = \frac{0}{0 + 100} = 0 \quad (0.5 \text{ point})$$

$$\epsilon_{2,R} = \frac{100}{200 + 100} \approx 0.33 \quad (0.5 \text{ point})$$

Entropy:

$$H_{1,L} = -\frac{150}{150 + 50} \ln\left(\frac{150}{150 + 50}\right) - \frac{50}{150 + 50} \ln\left(\frac{50}{150 + 50}\right) \approx 0.56 \quad (0.5 \text{ point})$$

$$H_{1,R} = -\frac{50}{150 + 50} \ln\left(\frac{50}{150 + 50}\right) - \frac{150}{150 + 50} \ln\left(\frac{150}{150 + 50}\right) \approx 0.56 \quad (0.5 \text{ point})$$

$$H_{2,L} = -\frac{0}{0 + 100} \ln\left(\frac{0}{0 + 100}\right) - \frac{100}{0 + 100} \ln\left(\frac{100}{0 + 100}\right) = 0 \quad (0.5 \text{ point})$$

$$H_{2,R} = -\frac{200}{200 + 100} \ln\left(\frac{200}{200 + 100}\right) - \frac{100}{200 + 100} \ln\left(\frac{100}{200 + 100}\right) \approx 0.64 \quad (0.5 \text{ point})$$

Gini impurity:

$$G_{1,L} = 1 - \left(\frac{150}{150 + 50}\right)^2 - \left(\frac{50}{150 + 50}\right)^2 = 0.375 \approx 0.38 \quad (0.5 \text{ point})$$

$$G_{1,R} = 1 - \left(\frac{50}{150 + 50}\right)^2 - \left(\frac{150}{150 + 50}\right)^2 = 0.375 \approx 0.38 \quad (0.5 \text{ point})$$

$$G_{2,L} = 1 - \left(\frac{0}{0 + 100}\right)^2 - \left(\frac{100}{0 + 100}\right)^2 = 0 \quad (0.5 \text{ point})$$

$$G_{2,R} = 1 - \left(\frac{200}{200 + 100}\right)^2 - \left(\frac{100}{200 + 100}\right)^2 \approx 0.44 \quad (0.5 \text{ point})$$

1.2 (6 pts) Compare the quality of \mathcal{T}_1 and \mathcal{T}_2 (that is, the two different splits of the root) based on conditional entropy (base e), weighted Gini impurity and total classification error. Intuitively, which of \mathcal{T}_1 or \mathcal{T}_2 appears to be a better split to you (there may not necessarily be one correct answer to this)? Based on your conditional entropy, Gini

impurity and classification error calculations, which of the metrics appear to be more suitable choices to decide which variable to split on?

Let $p_1 = \frac{150+50}{400} = 0.5$ be the fraction of examples that belong to left leaf of \mathcal{T}_1 , and $p_2 = \frac{0+100}{400} = 0.25$ be the fraction of examples that belong to left leaf of \mathcal{T}_2 . Then the total classification error for \mathcal{T}_1 and \mathcal{T}_2 are respectively:

$$\epsilon_1 = p_1\epsilon_{1,L} + (1 - p_1)\epsilon_{1,R} = 0.25 \quad (0.5 \text{ point})$$

$$\epsilon_2 = p_2\epsilon_{2,L} + (1 - p_2)\epsilon_{2,R} = 0.25 \quad (0.5 \text{ point})$$

So they are as good in terms of classification error. (0.5 points)

The conditional entropy for \mathcal{T}_1 and \mathcal{T}_2 are respectively:

$$\epsilon_1 = p_1H_{1,L} + (1 - p_1)H_{1,R} \approx 0.56 \quad (0.5 \text{ point})$$

$$\epsilon_2 = p_2H_{2,L} + (1 - p_2)H_{2,R} \approx 0.48 \quad (0.5 \text{ point})$$

So \mathcal{T}_2 is better in terms of conditional entropy. (0.5 points)

The weighted Gini impurity for \mathcal{T}_1 and \mathcal{T}_2 are respectively:

$$\epsilon_1 = p_1G_{1,L} + (1 - p_1)G_{1,R} \approx 0.38 \quad (0.5 \text{ point})$$

$$\epsilon_2 = p_2G_{2,L} + (1 - p_2)G_{2,R} \approx 0.33 \quad (0.5 \text{ point})$$

So \mathcal{T}_2 is also better in terms of weighted Gini impurity. (0.5 points)

Since \mathcal{T}_2 leads to a pure (100% certain) node, it is probably preferable to \mathcal{T}_1 . Therefore conditional entropy and Gini impurity appear to be more suitable measures to decide which variable to split on. They are more sensitive to changes in the node probabilities than classification error. (1.5 points)

Problem 2: Gaussian Mixture Model and EM (10pts+5pts Bonus)

In the lecture we applied EM to learn Gaussian Mixture Models (GMMs) and showed the M-Step without a proof. Now it is time that you prove it.

Consider a GMM with the following PDF of \mathbf{x}_n :

$$p(\mathbf{x}_n) = \sum_{k=1}^K \omega_k (\mathbf{x}_n \mid \mu_k, \Sigma_k) = \sum_{k=1}^K \frac{\omega_k}{(\sqrt{2\pi})^D |\Sigma_k|^{\frac{1}{2}}} \exp \left(-\frac{1}{2} (\mathbf{x}_n - \mu_k)^T \Sigma_k^{-1} (\mathbf{x}_n - \mu_k) \right)$$

where $K \in \mathbb{N}$ is the number of Gaussian components, $D \in \mathbb{N}$ is dimension of a data point \mathbf{x}_n . This GMM has K tuples of model parameters $(\mu_k, \Sigma_k, \omega_k)$, which stand for the mean vector, covariance matrix, and component weight of the k -th Gaussian component. $|\Sigma|$ denotes the determinant of matrix Σ . For simplicity, we further assume that all components are isotropic Gaussian, i.e., $\Sigma_k = \sigma_k^2 I$.

2.1 (10 pts) Find the MLE of the expected complete log-likelihood. Equivalently, find the optimal solution to the following optimization problem.

$$\begin{aligned} \underset{\omega_k, \mu_k, \Sigma_k}{\operatorname{argmax}} \quad & \sum_n \sum_k \gamma_{nk} \ln \omega_k + \sum_n \sum_k \gamma_{nk} \ln (\mathbf{x}_n \mid \mu_k, \Sigma_k) \\ \text{s.t.} \quad & \omega_k \geq 0 \\ & \sum_{k=1}^K \omega_k = 1 \end{aligned}$$

where γ_{nk} is the posterior of latent variables computed from the E-Step.

Hint: Given $a_1, \dots, a_K \in \mathbb{R}^+$, the solution to the following optimization problem over q_1, \dots, q_K :

$$\begin{aligned} \underset{q_k}{\operatorname{argmax}} \quad & \sum_{k=1}^K a_k \ln q_k, \\ \text{s.t.} \quad & q_k \geq 0, \\ & \sum_{k=1}^K q_k = 1. \end{aligned}$$

is given by:

$$q_k^* = \frac{a_k}{\sum_{k'} a_{k'}}$$

To find $\omega_1, \dots, \omega_K$, we simply solve

$$\begin{aligned} \underset{\omega}{\operatorname{argmax}} \quad & \sum_n \sum_k \gamma_{nk} \ln \omega_k. \\ \text{s.t.} \quad & \omega_k \geq 0 \\ & \sum_{k=1}^K \omega_k = 1 \end{aligned} \quad (2 \text{ points})$$

The solution is

$$\omega_k^* = \frac{\sum_n \gamma_{nk}}{\sum_k \sum_n \gamma_{nk}} = \frac{\sum_n \gamma_{nk}}{\sum_n 1} = \frac{\sum_n \gamma_{nk}}{N}. \quad (1 \text{ points})$$

To find μ_k and σ_k , we solve for each k

$$\begin{aligned} \operatorname{argmax}_{\mu_k, \Sigma_k} \sum_n \gamma_{nk} \ln(\mathbf{x}_n | \mu_k, \Sigma_k) &= \operatorname{argmax}_{\mu_k, \Sigma_k} \sum_n \gamma_{nk} \ln \left[\frac{1}{(\sqrt{2\pi}\sigma_k)^D} \exp \left(-\frac{1}{2\sigma_k^2} \|\mathbf{x}_n - \mu_k\|^2 \right) \right] \\ &= \operatorname{argmax}_{\mu_k, \Sigma_k} \sum_n \gamma_{nk} \left(-D \ln \sigma_k - \frac{\|\mathbf{x}_n - \mu_k\|^2}{2\sigma_k^2} \right), \end{aligned}$$

where D is the length of \mathbf{x}_n .

(3 points)

First we set the derivative w.r.t. μ_k to 0:

$$\frac{1}{\sigma_k^2} \sum_n \gamma_{nk} (\mathbf{x}_n - \mu_k) = 0,$$

which gives

$$\mu_k^* = \frac{\sum_n \gamma_{nk} \mathbf{x}_n}{\sum_n \gamma_{nk}} \quad (2 \text{ points})$$

Next we set the derivative w.r.t. σ_k to 0:

$$\sum_n \gamma_{nk} \left(-\frac{D}{\sigma_k} + \frac{\|\mathbf{x}_n - \mu_k\|^2}{\sigma_k^3} \right) = 0.$$

Solving for σ_k gives

$$(\sigma_k^*)^2 = \frac{\sum_n \gamma_{nk} \|\mathbf{x}_n - \mu_k^*\|^2}{D \sum_n \gamma_{nk}}. \quad (2 \text{ points})$$

2.2 (Bonus) (5 pts) The posterior probability of z in GMM can be seen as a *soft* assignment to the clusters; In contrast, K-means assign each data point to one cluster at each iteration (*hard* assignment). Describe how to set the model parameters such that the GMM in the previous question reduces to a K-means; please also derive $p(z_n = k | \mathbf{x}_n)$ in this case.

Set all $\sigma_k = \sigma \rightarrow 0$ and $\omega_k = \frac{1}{K}$, we have

$$p(\mathbf{x}_n, z_n = k) \propto \exp -\frac{1}{2\sigma^2} \|\mathbf{x}_n - \mu_k\|^2,$$

where constant terms are ignored.

(2 points)

The posterior then becomes

$$p(z_n = k | \mathbf{x}_n) = \lim_{\sigma \rightarrow 0} \frac{\exp\{-\frac{1}{2\sigma^2} \|\mathbf{x}_n - \mu_k\|^2\}}{\sum_j \exp\{-\frac{1}{2\sigma^2} \|\mathbf{x}_n - \mu_j\|^2\}} \rightarrow \begin{cases} 1, & \text{if } k = \arg_j \min \|\mathbf{x}_n - \mu_j\|^2 \\ 0, & \text{otherwise.} \end{cases}$$

(3 points)

Programming-based Questions

As in previous homeworks, you need to have your coding environment setup for this part. We use python3 (version ≥ 3.7) in our programming-based questions. There are multiple ways you can install python3, for example:

- You can use **conda** to configure a python3 environment for all programming assignments.
- Alternatively, you can also use **virtualenv** to configure a python3 environment for all programming assignments

After you have a python3 environment, you will need to install the following python packages:

- numpy
- matplotlib (for plotting figures)

Note: You are **not allowed** to use other packages such as *tensorflow*, *pytorch*, *keras*, *scikit-learn*, *scipy*, etc. for 3.1-3.2. If you have other package requests, please ask first before using them. You are **allowed** to use any packages for 3.3-3.5.

Download the files for the programming part from <https://vatsalsharan.github.io/fall22/hw3.zip>.

Problem 3: Exploring Decision Trees and Random Forests (12pts)

In this question, we will observe the effect of different hyperparameters in training decision trees and random forests, and also visualize what features the random forest model is using to make predictions. We will do this on a Colab notebook HW4-Exploring-Random-Forests.ipynb.

Instructions to run the notebook: Upload this file to your USC Google Drive. Then, add Google Colab to your Google App by New \rightarrow More \rightarrow Connect more apps \rightarrow Type in Google Colab and install it, as in Fig. 1. After that, you can run the notebook with your browser.

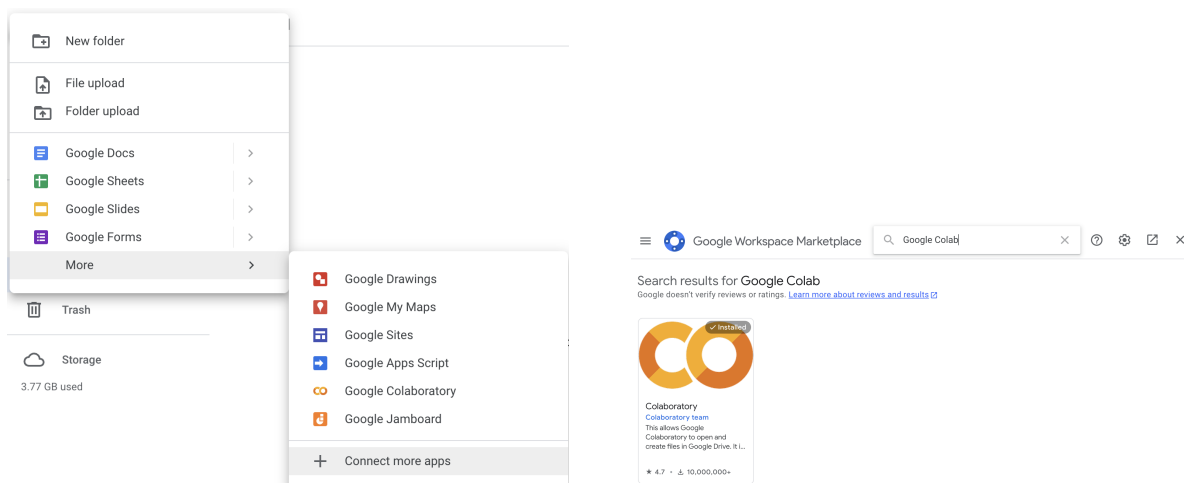


Figure 1: Screenshots showing how to install Colab.

We use a variant of the MNIST dataset for this problem. As mentioned in HW 3, the MNIST dataset contains images of handwritten digits (0, 1, 2, ..., 9) and is generally used for the (10) digit classification problem. Here, we work with a binary classification task of predicting whether the digit is less than 5 or not. Fig. 2 shows some samples images from the dataset with original and binary labels, respectively.

You should go through the code we provide and understand what's happening, but for the purpose of answering the questions you will mainly have to run the code and understand the results. All the models (decision trees, random forests, etc.) are imported from the *sklearn* library. You should feel free to explore the role of other hyperparameters and other methods (such as bagging, boosting), and go through the documentation to better understand these things, but for this question, we will focus on decision trees and random forests.




Figure 2: Some samples images from the MNIST dataset with original (left) and binary (right) labels, respectively.

We will look at the effect of 5 parameters:

- `max_depth`: The maximum depth of the decision tree.
- `n_estimators`: The number of decision trees in the forest.
- `min_samples_leaf`: The minimum number of samples required to be a leaf node.
- `max_samples`: The number of samples to draw from the training set to train each decision tree in the forest.
- `max_features`: The number of features to consider when looking for the best split for any node.

To observe the effect of a parameter, we look at train and test accuracy for different values of that parameter while keeping the rest of the parameters fixed.

3.1 (2 pts) The first set of plots shows the train and test accuracy for different values of `max_depth` using a decision tree and a random forest, respectively. How do the accuracies and the generalization gap vary with `max_depth` in these cases? For a particular value of `max_depth`, how does the generalization gap for the decision tree compare with that of the random forest? Explain your observations.

In both the cases, as `max_depth` increases, the train accuracy keeps increasing, while the test accuracy increases upto a point and then saturates. The generalization gap also keeps increasing as `max_depth` is increased. (1pt)

For a fixed value of `max_depth`, the generalization gap for the decision tree is much higher than the random forest. This is because a single decision tree is trained using all the features and its predictions depend on particular set of features, whereas the random forest leverages multiple decision trees to make its predictions, each of which is trained using only a subset of the features. This helps avoid the risk of overfitting onto a specific set of features. (1pt)

3.2 (2 pts) The next plot shows the train and test accuracy for different values of `n_estimators` for a random forest. Comment on your observations regarding the accuracies and the generalization gap. What value(s) (range or approximate values are enough) would you prefer to use for this parameter? Give reasons why. Hint: Are there any drawbacks to using very high values of `n_estimators`?

As `n_estimators` increases, the train and test accuracies increase upto a point and then saturate. The generalization gap remains similar throughout. (1pt)

As the accuracies saturate around `n_estimators= 50`, it seems like values in the range 50 – 100 would work well. Although using higher values of `n_estimators` would give similar performance, it comes at the cost of increased training time. (1pt)

For the next three parts of this question, we will also look at how the size of the training set influences the accuracy trends for a given parameter. In each case, for the first plot, the training set consists of 4000 samples whereas for the second plot, it contains 1000 samples.

3.3 (2 pts) The next set of plots shows the train and test accuracy for different values of `min_samples_leaf` for a random forest. Taking into account the behaviours for different training set sizes, explain your observations for very low and very high values of `min_samples_leaf`. What do you conclude from this trend? What could be the reasons for such a behaviour?

In both the cases, for low values of `min_samples_leaf`, the train/test accuracies as well as the generalization gap is high, whereas for high values of `min_samples_leaf`, the train/test accuracies as well as the generalization gap is low. (1pt)

There is a tradeoff between high train/test accuracy and low generalization gap in both the cases. Very low values of `min_samples_leaf` can make the model prone to overfitting on the training data, whereas for very high values, the trees might be too shallow and thus, not learn a very complex decision boundary. (1pt)

3.4 (2 pts) The next set of plots shows the train and test accuracy for different values of `max_samples` for a random forest. What do you observe for very low and very high values of `max_samples`? Would you prefer to use low, intermediate or high values for this parameter in both the cases? Give reasons why. Hint: How does the size of the training set influence the choice of this parameter?

In both the cases, for low values of `max_samples`, the train/test accuracies as well as the generalization gap is low, whereas for high values of `max_samples`, the train/test accuracies as well as the generalization gap is high. (1pt)

In the first case, both intermediate and high values seem fine, while in the second case, it is better to use intermediate values. In the first case, even if all the trees in the forest rely on all the samples, it doesn't hurt the accuracy as the training set is larger, so there are less chances of overfitting. However, in the second case, the smaller training set size makes the setting prone to overfitting when the trees use all the samples. Using different subsets of the training set for different trees helps avoid that in this case. (1pt)

3.5 (2 pts) The next set of plots shows the train and test accuracy for different values of `max_features` for a random forest. Comment on your observations regarding the accuracies and the generalization gap for the two training set sizes. What is the best range of values for this parameter in both the cases? Is it similar/different? Explain your observations.

In the first plot, as `max_features` increases, both train and test accuracies increase upto a point, fluctuate a little, and then saturate. The generalization gap seems to be increasing (at a slow rate) throughout. In the second plot, as `max_features` increases, both train and test accuracies increase upto a point, fluctuate a little, and then start decreasing rapidly. The generalization gap increases throughout. (1pt)

In the first case, any value greater 25 seems fine, while in the second case, values in the range 25 – 100 seem to work well. In the first case, relying on a very large number of features doesn't seem to hurt as much as in the second case. In the first case, the training set is larger, so the sample distribution of the features is closer to the population distribution than the second case. Thus, the second case is more prone to overfitting when using a large value of `max_features`. (1pt)

3.6 (2 pts) In class, we discussed how ensembles are usually not as interpretable as a single decision tree. While this is true, there are still ways to explore which features are used the most by our ensemble. We will explore one such technique in this part.

We visualize the *feature importances* of a random forest model trained for the binary classification task on the MNIST dataset. Intuitively, features with higher importance are the pixels which are used more often in the decision trees in the forest and which lead to better splits, i.e. which contribute more in improving the performance of the model. For more details, you can see Section 18.6.1 of the PML book. The last plot shows the importances of different pixels/portions of the image for a trained random forest model to make its predictions. What portions of the image does the model seem to be focusing on? In other words, can you think of reasons why the pixels with higher importance are indeed important for the prediction task (classifying whether the digit is smaller than 5 or not)? As is usually true for such open-ended questions, there can be multiple correct answers here and we're looking more for your reasoning than a specific answer.

The model seems to be focusing on portions of the image which are common in digits 5, 6, 8 and 9. As the prediction task is classifying whether the digit is smaller than 5 or not, this seems expected as the digits 0-4 have less common portions with this image. Hence, focusing on these portions allows the model to distinguish between these two classes. (2pts)

Problem 4: PCA for Learning Word Embeddings (35pts+10pts Bonus)

This question is about *word embeddings*. A word embedding is simply a vector space representation of words which captures some of the semantic and syntactic structures in the language—for example, words similar in meaning have representations which are close to each other in the vector space. Word embeddings have taken NLP by storm in the last few years, and have become the backbone for numerous NLP tasks such as question answering or machine translation. There are neural approaches to learning word embeddings, but in this question we will study a simple PCA-based scheme which does a surprisingly good job at learning word embeddings.

We have created a word co-occurrence matrix \mathbf{M} of the 10000 most frequent words from a Wikipedia corpus with 1.5 billion words. The co-occurrences were obtained by using a sliding window of length 5 across the Wikipedia corpus. Entry M_{ij} of the matrix denotes the number of times words i and j occur in the corpus within the same sliding window. The file `co_occur.csv` contains the symmetric co-occurrence matrix. `dictionary.txt` contains the dictionary for interpreting this matrix, the i th row of the dictionary is the word corresponding to the i th row or column of \mathbf{M} . The dictionary is sorted according to the word frequencies. Therefore the first word in the dictionary—"the" is the most common word in the corpus and the first row or column of \mathbf{M} contains the co-occurrence counts of "the" with every other word in the dictionary.

We provide some starter code in the file `hw4-pca.py` with some useful functions (e.g. to read files, generate plots, etc.) which can be used directly, and instructions on how to complete the functions required for this problem.

4.1 (6 pts) First, read the co-occurrence matrix and the list of all words from the given files. Let the matrix \mathbf{M} be the 10000×10000 matrix of word co-occurrences. In light of the power law distribution of word occurrences, it is more suitable to work with the normalized matrix $\tilde{\mathbf{M}}$ such that each entry $\tilde{M}_{ij} = \log(1 + M_{ij})$. Now, apply PCA on $\tilde{\mathbf{M}}$ to find its rank-100 approximation \mathbf{VDV}^T . Plot the eigenvalues $\text{diag}(\mathbf{D})$. Does the matrix seem to be low rank? Why?

The graph shows that after the first eigenvalue, there is an enormous drop in the value of the eigenvalues. The values keep decreasing and around 50 to 60, the graph starts to flatten out fully, which means the other ranks most likely do not add as much info and that \mathbf{M} might be a low-rank matrix.

Rubrics: 4 pts for the graph (3 pts for the plot and 1 pt for correct range of values on the y-axis), 2 pts for correct answer and reason.

4.2 (6 pts) In this question, we will use word embeddings to find word(s) which are 'similar' to a given word. First, normalize the rows of \mathbf{V} such that each row has unit ℓ_2 norm. We will regard the i th row of \mathbf{V} as the embedding of the i th word. Next, we will define a similarity metric for the word embeddings. We will use the cosine-similarity as the similarity metric. As all the vectors have unit ℓ_2 norm, the cosine similarity between two words i and j with embeddings \mathbf{w}_i and \mathbf{w}_j is equal to the inner product $\langle \mathbf{w}_i, \mathbf{w}_j \rangle$. Now that we have a similarity metric defined, we can




Figure 3: Plot showing the first 100 eigenvalues.

have some fun with these embeddings by querying for the closest word to any word we like! Try finding the closest words to some common words, such as “learning”, “university”, “california”, and comment on your observations.

The closest words to “learning”, “university”, “california” were “teaching”, “college”, “florida”, respectively. It seems like words that seem related to each other have high similarity in the embedding space that we get using the PCA approximation.

Rubrics: 4 pts for getting similar words, 2 pts for observations.

4.3 (9 pts) We’ll now interpret the principal components/eigenvectors (columns of \mathbf{V}). For any i , denote v_i as the eigenvector corresponding to the i th largest eigenvalue. Note that the entries of this vector correspond to the 10000 words in our dictionary, we’ll call these our 10000 variables. By sorting the entries of v_i by absolute value, and observing what the top 10 variables and their (signed) entries are, we can infer what information the i th eigenvector roughly captures. Can you find 8 interesting eigenvectors, and point out what semantic or syntactic structures they capture? Can you do this for all eigenvectors? Hint: What do you observe about PCs or eigenvectors with small eigenvalues?

Sample PCs and top 10 words for each of them:

PC 0 gives: ['marshall' 'harvey' 'curtis' 'spencer' 'oliver' 'mason' 'lloyd' 'owen' 'howard' 'warren']
 PC 1 gives: ['oliver' 'curtis' 'howard' 'gordon' 'barry' 'anthony' 'harvey' 'lloyd' 'spencer' 'leslie']
 PC 5 gives: ['went' 'put' 'took' 'allowed' 'turned' 'came' 'saw' 'texts' 'led' 'take']
 PC 6 gives: ['ancient' 'inspired' 'famous' 'wade' 'medieval' 'bradley' 'greek' 'contemporary' 'ryan' 'culture']
 PC 7 gives: ['games' 'teams' 'tournament' 'team' 'contest' 'tour' 'competition' 'clubs' 'league' 'championship']
 PC 8 gives: ['competed' 'standings' 'scorer' 'finals' 'sv' 'substituted' 'liga' 'warren' 'knockout' 'italia']
 PC 9 gives: ['commander' 'ii' 'tank' 'command' 'force' 'carrier' 'forces' 'military' 'naval' 'army']
 PC 21 gives: ['glass' 'walls' 'wall' 'clay' 'hanging' 'inside' 'panels' 'open' 'nominal' 'front']
 PC 22 gives: ['wales' 'cincinnati' 'scotland' 'ian' 'miami' 'milwaukee' 'stuart' 'colin' 'toledo' 'nigel']

It looks like PCs 0 and 1 capture male first names, PC 5 captures action verbs, PC 6 has words related to history, the ones in PCs 7 and 8 are related to sports and competition, PC 9 captures words related to military, while the ones in PC 21 are related to walls, and PC 22 captures names of places.

Not all PCs make sense because as the strongest PCs are covered, like the ones described above and vectors that

describe concepts like grammar or articles, more obscure concepts are left to be described so the remaining PCs, especially near the bottom, will not be as obvious to human eye but will still have some sort of concept described with the PCs, just ones we won't be able to observe easily or ones that aren't mentioned as often.

Rubrics: 4 pts for getting the top 10 words for 8 eigenvectors/PCs (0.5 each) (these can be different from the ones shown here, these are just examples), 4 pts for mentioning the semantic or syntactic structures (0.5 each), 1 pt for the answer and explanation.

4.4 (14 pts) In this question, we will explore a curious property of the word embeddings—that certain directions in the embedded space correspond to specific syntactic or semantic concepts. Let \mathbf{v}_1 be the word embedding for “woman” and \mathbf{v}_2 be the word embedding for “man”. Let $\mathbf{v} = \mathbf{v}_1 - \mathbf{v}_2$, and $\hat{\mathbf{v}} = \frac{\mathbf{v}}{\|\mathbf{v}\|}$.

4.4.1 (6 pts) Project the embeddings of the following words onto $\hat{\mathbf{v}}$: *boy, girl, brother, sister, king, queen, he, she, john, mary, wall, tree*. Present a plot of projections of the embeddings of these words marked on a line. For example, if the projection of the embedding for “girl” onto $\hat{\mathbf{v}}$ is 0.1, then you should label 0.1 on the line with “girl”. What do you observe?

It can be seen that words used primarily with males like “John” or “brother” are on the negative projection of $\hat{\mathbf{v}}$ and




Figure 4: Plot showing projections of words in part 4.4.1.

words like “queen” and “sister” have a positive projection onto $\hat{\mathbf{v}}$. This makes sense since $\mathbf{v} = \mathbf{v}_1(\text{woman}) - \mathbf{v}_2(\text{man})$ and this positive projection imply a positive correlation with feminine gender alignment and negative means less feminine and as a result more male.

Rubrics: 3 pts for correct plot, 2 pts for correct observation.

4.4.2 (8 pts) Present a similar plot of the projections of the embeddings of the following words onto $\hat{\mathbf{v}}$: *math, history, nurse, doctor, pilot, teacher, engineer, science, arts, literature, bob, alice*. What do you observe? Why do you think this is the case? Do you see a potential problem with this? Remember that word embeddings are extensively used across NLP. Suppose LinkedIn’s search feature to find suitable candidates for a job used such word embeddings to find candidates who best match a search term or job description. What might be the result of this?

It can be seen that words like “engineer” and “bob” and “math” have a less feminine gender alignment and more of a male one, and words like “nurse” or “literature” have a more feminine alignment, since it has a positive projection onto $\hat{\mathbf{v}}$.

This could be the case because there is a gender bias within occupations from the Wikipedia corpus that people as-




Figure 5: Plot showing projections of words in part 4.4.2.

sociate more technical roles with males and more artistic roles with females, so in other words the language shows a gender bias.

This is a bad idea for something like linkedin because depending on language that shows a gender bias in occupations will alienate people so that they are associated with their respective gendered roles. So male nurses would get marginalized since the embedding for "nurse" has more of a female orientation and female engineers would also get marginalized on search queries because "engineer" has more of a male orientation.

Rubrics: 2 pts for correct plot, 1.5 pts for correct observation, 3.5 pts for correct answers and explanation (1.5+2).

4.5 (Bonus) (10 pts) In this question we will explore in more depth the property that directions in the embedded space correspond to semantic or syntactic concepts.

Because word embeddings capture semantic and syntactic concepts, they can be used to solve word analogy tasks. For example, consider an analogy question— "*man is to woman as king is to ____*", where the goal is to fill in the blank space. This can be solved by finding the word whose embedding is closest to $\mathbf{w}_{\text{woman}} - \mathbf{w}_{\text{man}} + \mathbf{w}_{\text{king}}$ in cosine similarity. You can do this by a nearest neighbor search across the entire dictionary—excluding the three words *man*, *woman*, *king* which already appear in the analogy as they cannot be valid answers to the question. Here \mathbf{w}_i represents the word embedding for the word i . Refer to Fig. 6 for why this makes sense in light of the fact that directions in the embedded space correspond to semantic/syntactic concepts.

We have provided a dataset, `analogy_task.txt`, which tests the ability of word embeddings to answer analogy questions, such as those represented in Fig. 6. Using the cosine similarity metric, find and report the accuracy of the word embeddings you have constructed on the word analogy task.

The accuracy obtained is around 55.33%.

Rubrics: 10 pts for correct implementation and correct accuracy. Partial credits to be awarded based on how much of the implementation is correct.

Deliverables: Discussions for all the parts. Plots for parts 4.1, 4.4.1, 4.4.2. Code for all the parts as a separate Python file `hw4-pca.py`.




Figure 6: Figure denoting how word embeddings might encode gender, tense or country-capital relationships. Image source.

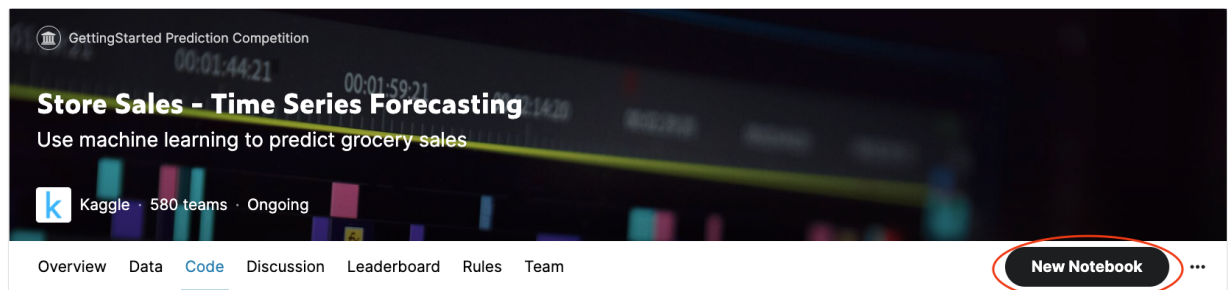
Problem 5: Project Preparation (16pts)

In this part, you will implement a simple linear regression model for the final project. At the same time, you will also explore the use of the Pandas package, a very popular Python library for data analysis. Here is a tutorial for the Pandas package. You should do this part together with your project team (each team will make one submission).

Early bird bonus We encourage you to start early on the project. The first 5 teams among all teams in CSCI-567 class at the **HW4 submission deadline** can get 10 bonus points. To get this bonus, at least one of the team members will have to give a 5-minute presentation on their approach in the discussion session on November 17. We will consider the final ranking on the public leaderboard as of **2pm PST on Nov 16** for this part of bonus points.

Setup We will complete this part of assignment on a jupyter notebook `project-starter.ipynb`. To load the notebook into Kaggle, follow these steps:

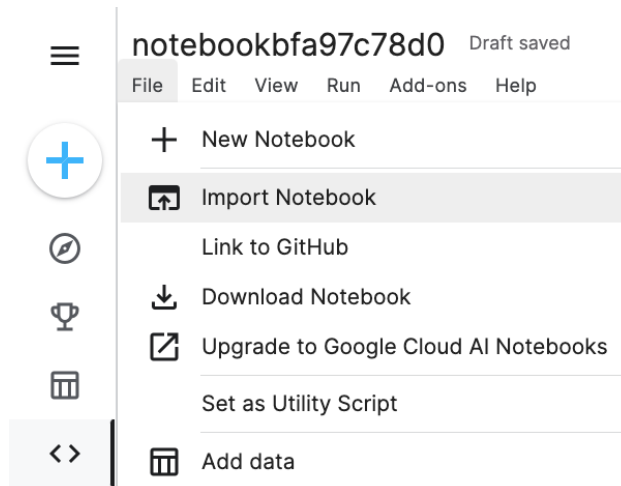
- In the “Code” tag of the competition, click button “New Notebook”.



Notebooks

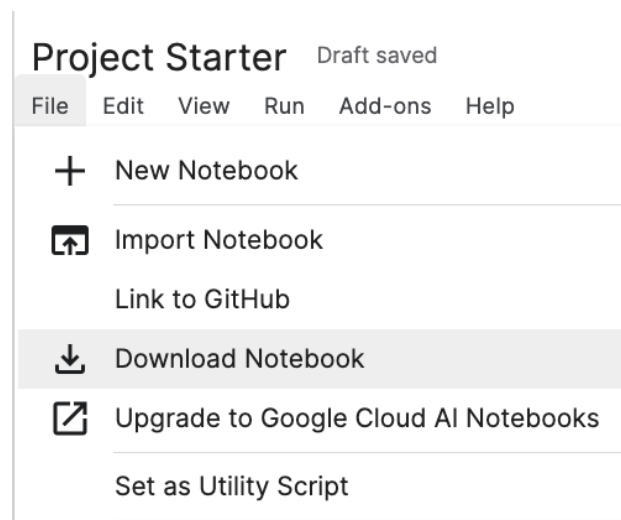


- Then, select “File/Import Notebook”, and import `regression.ipynb` into Kaggle.



Deliverables Notebook with completed code, output, and the discussion for 5.4.

Submission Instructions Each project team (not homework team) will make one submission for this problem (there will be a separate Gradescope link to submit solutions for this problem). Please submit your notebook *along with the output* for grading. If you are using Kaggle notebook, you can export the notebook by “File/Download Notebook”, see figure below:



General Tip 1: before getting started, please read the dataset description carefully. Make sure you understand the datatype of each entries and the meanings behind them.

General Tip 2: all of the tasks in this part can be done within three lines of code. You should spend most of your time reading documentations to understand certain methods provided by the Pandas package, instead of implementing the functionalities yourself.

Part I - Feature Engineering (12 pts)

One of the most important steps in developing a machine learning model is feature engineering. In this part, we will walk you through the process of constructing a simple set of features for the final project.

5.1 Moving average of oil prices (6 pts) In this task, you are asked to compute the moving average of the oil prices with window size 7.

5.1.1 (2 pts) In Section 1 of the notebook, after loading the oil price into `data_oil`, compute `data_oil["ma_oil"]` as the moving average of past seven recorded oil prices (note that the oil prices of some dates are missing, and we will skip them in the computation).

5.1.2 (4 pts) After computing `data_oil["ma_oil"]`, you should now deal with the missing dates in `data_oil["ma_oil"]`. A DataFrame `calendar` with contiguous dates from 2013-01-01 to 2017-08-31 is created for you. You are asked to complete the following:

- (a) Merge DataFrames `calendar` and `data_oil`. The merged DataFrame should have contiguous dates from 2013-01-01 to 2017-08-31 as indexes, and the same columns as `data_oil`, where missing values are denoted by `NaN`.
- (b) Fill in the missing values of `data_oil["ma_oil"]`. There are many ways to deal with missing values, such as fill it with some default value, fill it with the last valid observation, or fill it with the next valid observation. Here, you are asked to fill in each missing value by the last valid observation. Perform this operation *in place* on `data_oil["ma_oil"]`.

Compare `ma_oil` column in the output with the solution. Get points for 5.1.1 if the values are correct. Get points for 5.1.2 (a) if the dates are contiguous. Get points for 5.1.2 (b) if NaNs are properly filled.

5.2 Workday feature (3 pts) Next, we will create an extra feature indicating whether each date is a workday. Here, we only make use of the fact that Monday to Friday are usually workdays, and we ignore all the holidays information in `holidays_events.csv` (which is left for you to explore). In Section 2 of the notebook, your task is to create a column `calendar["wd"]` such that a date between Monday and Friday has value `True`, and a date that is Saturday or Sunday has value `False`. Compare `wd` column in the output with the solution. Column `dofw` is not necessary.

5.3 Create trend feature (3 pts) Now we are ready to load all training and testing data into the notebook (Section 3). The last step before training our model is to create the trend feature: note that a linear regression model cannot read the dates, and we need to map the dates into certain features. Here, you are asked to implement the simplest trend feature: given a range of dates, map them into a range of integers starting from 1. You need to create a DataFrame `X` as the trend feature, which is then extended with the moving average of oil prices and workday indicators to form the final training features. Compare `trend` column in the output with the solution.

Part II - Train Your Model (4 pts)

We are now ready to train the model! Here, we will train a simple linear regression model, which only takes a few seconds.

5.4 Display training error by the type of product (2 pts) Let's show the training error to get a sense of how well the model is fitting the training data. The model is making sales predictions for each store and type of product. You are asked to show the training error within each type of product. Which family of product does the model has the most difficulty in predicting the sales? Compare output with the solution.

5.5 Make predictions on the test set (2 pts) Now you can use your trained model to make predictions on the test set. To do this, you need to first create the features for the test set (denoted by `X_test`) similar to the training set. Fill in your steps for computing `X_test` in the notebook. Compare output with the solution.

5.6 Make a submission You are now ready to make a submission. Simply open the "Competitions" tag on the right, and click the submit button to make the submission. You should get a score around 0.45.