

# CSCI 567: Machine Learning

Vatsal Sharan  
Fall 2022

Lecture 10, Nov 10

# Administrivia

- Start on your project if you haven't already!
  - Make groups (of 4) by tomorrow (Nov 11), minimum team size is 3.
  - Top teams as of Nov 16 can get a bonus!
- HW4 is due in about one weeks (Nov 16 at 2pm).
  - We'll release another question on Gaussian mixture models tomorrow.
- Today's plan:
  - Clustering
  - Gaussian Mixture Models and Expectation Maximization (EM)
  - In the discussion, we will go over popular evaluation metrics for supervised learning

# A simplistic taxonomy of ML

## **Supervised learning:**

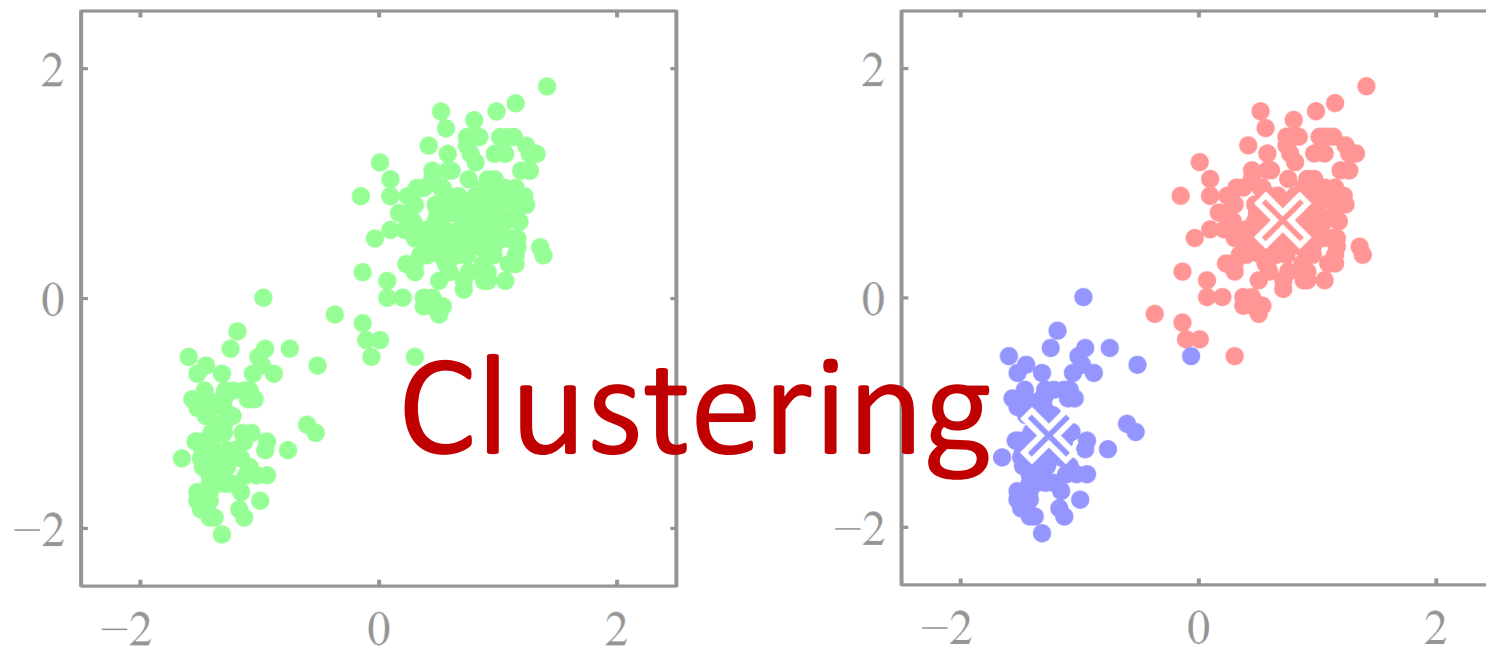
Aim to predict  
outputs of future  
datapoints

## **Unsupervised learning:**

Aim to discover  
hidden patterns and  
explore data

## **Reinforcement learning:**

Aim to make  
sequential decisions



# Clustering



- Introduction
- Formalizing and solving the objective (alternating minimization)
- $k$ -means algorithm

# Clustering: Informal definition

**Given:** a set of data points (feature vectors), *without labels*

**Output:** group the data into some clusters, which means

- **assign** each point to a specific cluster
- find the **center** (representative/prototype/...) of each cluster





# Clustering: More formal definition

**Given:** data points  $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^d$  and #clusters  $k$  we want

**Output:** group the data into  $k$  clusters, which means,

- find **assignment**  $\gamma_{ij} \in \{0, 1\}$  for each data point  $i \in [n]$  and  $j \in [k]$  s.t.  $\sum_{j \in [k]} \gamma_{ij} = 1$  for any fixed  $i$
- find the cluster **centers**  $\mu_1, \dots, \mu_k \in \mathbb{R}^d$



# Many applications

Clustering is one of the most fundamental ML tasks, with many applications:

- recognize communities in a social network
- group similar customers in market research
- image segmentation
- accelerate other algorithms (e.g. nearest neighbor classification)
- ...



# Clustering

- Introduction
- Formalizing and solving the objective (alternating minimization)
- $k$ -means algorithm

# Formal objective

As with PCA, no ground-truth to even measure the quality of the answer (*no labels given*).

What is the high-level goal here?

We want to partition the points into  $k$  clusters, such that points within each cluster are close to their cluster center.

We can turn this into an optimization problem, find  $\gamma_{ij}$  and  $\mu_j$  to minimize

$$F(\{\gamma_{ij}\}, \{\mu_j\}) = \sum_{i=1}^n \sum_{j=1}^k \gamma_{ij} \|\mathbf{x}_i - \mu_j\|_2^2$$

i.e. the **sum of squared distances of each point to its center**. This is the “ **$k$ -means” objective**.

# How to solve this? Alternating minimization

Unfortunately, finding the exact minimizer of the  $k$ -means objective is *NP-hard!*

Therefore, we use a heuristic (*alternating minimization*) that alternately minimizes over  $\{\gamma_{ij}\}$  and  $\{\mu_j\}$ :

Initialize  $\{\mu_j^{(1)} : j \in [k]\}$

For  $t = 1, 2, \dots$

- find

$$\{\gamma_{ij}^{(t+1)}\} = \operatorname{argmin}_{\{\gamma_{ij}\}} F\left(\{\gamma_{ij}\}, \{\mu_j^{(t)}\}\right)$$

- find

$$\{\mu_j^{(t+1)}\} = \operatorname{argmin}_{\{\mu_j\}} F\left(\{\gamma_{ij}^{(t+1)}\}, \{\mu_j\}\right)$$

# Alternating minimization: Closer look

The first step

$$\begin{aligned}\min_{\{\gamma_{ij}\}} F(\{\gamma_{ij}\}, \{\boldsymbol{\mu}_j\}) &= \min_{\{\gamma_{ij}\}} \sum_i \sum_j \gamma_{ij} \|\mathbf{x}_i - \boldsymbol{\mu}_j\|_2^2 \\ &= \sum_i \min_{\{\gamma_{ij}\}} \sum_j \gamma_{ij} \|\mathbf{x}_i - \boldsymbol{\mu}_j\|_2^2\end{aligned}$$

is simply to **assign each  $x_i$  to the closest  $\mu_j$** , i.e.

$$\gamma_{ij} = \mathbb{I} \left[ j == \underset{c}{\operatorname{argmin}} \|\mathbf{x}_i - \boldsymbol{\mu}_c\|_2^2 \right]$$

for all  $j \in [k]$  and  $i \in [n]$ .

# Alternating minimization: Closer look

The second step

$$\begin{aligned}\min_{\{\boldsymbol{\mu}_j\}} F(\{\gamma_{ij}\}, \{\boldsymbol{\mu}_j\}) &= \min_{\{\boldsymbol{\mu}_j\}} \sum_i \sum_j \gamma_{ij} \|\mathbf{x}_i - \boldsymbol{\mu}_j\|_2^2 \\ &= \sum_j \min_{\boldsymbol{\mu}_j} \sum_{i:\gamma_{ij}=1} \|\mathbf{x}_i - \boldsymbol{\mu}_j\|_2^2\end{aligned}$$

is simply **to average the points of each cluster** (hence the name)

$$\boldsymbol{\mu}_j = \frac{\sum_{i:\gamma_{ij}=1} \mathbf{x}_i}{|\{i : \gamma_{ij} = 1\}|} = \frac{\sum_i \gamma_{ij} \mathbf{x}_i}{\sum_i \gamma_{ij}}$$

for each  $j \in [k]$ .

# Clustering

- Introduction
- Formalizing and solving the objective (alternating minimization)
- $k$ -means algorithm

# *k*-means algorithm

**Step 0** Initialize  $\mu_1, \dots, \mu_k$

**Step 1** For the centers  $\mu_1, \dots, \mu_k$  being fixed, assign each point to the closest center:










$$\gamma_{ij} = \mathbb{I} \left[ j == \underset{c}{\operatorname{argmin}} \|\mathbf{x}_i - \mu_c\|_2^2 \right]$$

**Step 2** For the assignments  $\{\gamma_{ij}\}$  being fixed, update the centers

$$\mu_j = \frac{\sum_i \gamma_{ij} \mathbf{x}_i}{\sum_i \gamma_{ij}}$$

**Step 3** Return to Step 1 if not converged (convergence means that all the assignments  $\gamma_{ij}$  are unchanged in Step 1).

# *k*-means algorithm: Example





# $k$ -means algorithm: Convergence

$k$ -means will converge in a finite number of iterations, why?

- objective strictly decreases at each step if the algorithm has not converged.

Why? For  $t = 1, 2, \dots$

- find

$$\begin{aligned}\{\gamma_{ij}^{(t+1)}\} &= \operatorname{argmin}_{\{\gamma_{ij}\}} F\left(\{\gamma_{ij}\}, \{\boldsymbol{\mu}_j^{(t)}\}\right) \\ &= \mathbb{I} \left[ j == \operatorname{argmin}_c \|\mathbf{x}_i - \boldsymbol{\mu}_c\|_2^2 \right]\end{aligned}$$

- find

$$\begin{aligned}\{\boldsymbol{\mu}_j^{(t+1)}\} &= \operatorname{argmin}_{\{\boldsymbol{\mu}_j\}} F\left(\{\gamma_{ij}^{(t+1)}\}, \{\boldsymbol{\mu}_j\}\right) \\ &= \frac{\sum_i \gamma_{ij} \mathbf{x}_i}{\sum_i \gamma_{ij}}\end{aligned}$$

# *k*-means algorithm: Convergence

*k*-means will converge in a finite number of iterations, why?

- objective strictly decreases at each step if the algorithm has not converged.
- #possible\_assignments is finite ( $k^n$ , exponentially large though)

Therefore, the algorithm must converge in at most  $k^n$  steps.

Why? More specifically, why can't the algorithm cycle between different clusterings?

- Suppose the algorithm finds the same clustering at time steps  $t_1$  and  $t_2$ .
- Since the objective function value decreases at every step, this means the same clustering (at time steps  $t_1$  and  $t_2$ ) has two different costs, which is not possible.
- Therefore, by contradiction, the algorithm cannot cycle between clusterings.

# *k*-means algorithm: Convergence

*k*-means will converge in a finite number of iterations, why?

- objective strictly decreases at each step if the algorithm has not converged.
- #possible\_assignments is finite ( $k^n$ , exponentially large though)

However

- it could take *exponentially many iterations* to converge
- and it *might not converge to the global minimum* of the *k*-means objective

# *k*-means algorithm: How to initialize?


There are different ways to initialize:

- randomly pick  $k$  points as initial centers  $\mu_1, \dots, \mu_k$
- or randomly assign each point to a cluster, then average to find centers
- or more sophisticated approaches (e.g. *k*-means++)

Initialization matters for **convergence**.

# *k*-means algorithm: Local vs Global minima

Simple example: 4 data points, 2 clusters, 2 different initializations




K-means converges immediately in both cases, but

- left has K-means objective  $L^2 = 4W^2$
- right has K-means objective  $W^2$ , *4 times better than left!*
- in fact, left is **local minimum**, and right is **global minimum**.

# *k*-means algorithm: Summary

- Clustering is a fundamental unsupervised learning task.
- *k*-means is an alternating minimization algorithm for the *k*-means objective.
- The algorithm always converges, but it can converge to a local minimum.
- Initialization matters a lot for the convergence. There are principled initialization schemes, which have guarantees on the solution they find (e.g. *k*-means++).



# Gaussian Mixture Model

- Introduction
- Learning the parameters
- EM algorithm
- EM for the Gaussian Mixture Model



# Gaussian mixture models

Gaussian mixture models (GMM) is a probabilistic approach for clustering

- more explanatory than minimizing the  $k$ -means objective
- can be seen as a soft version of  $k$ -means

To solve GMM, we will introduce a powerful method for learning probabilistic models: the **Expectation Maximization (EM) algorithm**.

# A generative model


For classification, we discussed the sigmoid model to “explain” how the labels are generated.

Similarly, for clustering, we want to come up with a probabilistic model  $p$  to “**explain**” **how the data is generated**.

That is, each point is **an independent sample** of  $\mathbf{x} \sim p$ .

Why do generative modelling?

- can generate data from  $p$
- can estimate probability of seeing any datapoint (useful for many tasks, such as for finding outliers/anomalies in data)



*What probabilistic model generates data like this?*

# GMM: Intuition

GMM is a natural model to explain such data.

Assume there are 3 ground-truth Gaussian models.

To generate a point, we

- first **randomly pick one of the Gaussian models**,
- then **draw a point according this Gaussian**.

Hence the name “**Gaussian mixture model**”.





Figure from Wikipedia

# GMM: Formal definition

A GMM has the following density function:

$$p(\mathbf{x}) = \sum_{j=1}^k \pi_j N(\mathbf{x} \mid \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)$$

where

- $k$ : the number of **Gaussian components** (same as #clusters we want)
- $\pi_1, \dots, \pi_k$ : **mixture weights**, a distribution over  $k$  components
- $\boldsymbol{\mu}_j$  and  $\boldsymbol{\Sigma}_j$ : **mean and covariance matrix** of the  $k$ -th Gaussian
- $N$ : the density function for a Gaussian

## Another view


By introducing a **latent variable**  $z \in [k]$ , which indicates cluster membership, we can see  $p$  as a **marginal distribution**

$$p(\mathbf{x}) = \sum_{j=1}^k p(\mathbf{x}, z = j) = \sum_{j=1}^k p(z = j)p(\mathbf{x}|z = j) = \sum_{j=1}^k \pi_j N(\mathbf{x} \mid \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)$$

$\mathbf{x}$  and  $z$  are both random variables drawn from the model

- $\mathbf{x}$  is **observed**
- $z$  is **unobserved/latent**

# An example




The conditional distributions are

$$p(\mathbf{x} \mid z = \text{red}) = N(\mathbf{x} \mid \boldsymbol{\mu}_1, \boldsymbol{\Sigma}_1)$$

$$p(\mathbf{x} \mid z = \text{blue}) = N(\mathbf{x} \mid \boldsymbol{\mu}_2, \boldsymbol{\Sigma}_2)$$

$$p(\mathbf{x} \mid z = \text{green}) = N(\mathbf{x} \mid \boldsymbol{\mu}_3, \boldsymbol{\Sigma}_3)$$



The marginal distribution is

$$\begin{aligned} p(\mathbf{x}) = & p(\text{red})N(\mathbf{x} \mid \boldsymbol{\mu}_1, \boldsymbol{\Sigma}_1) + p(\text{blue})N(\mathbf{x} \mid \boldsymbol{\mu}_2, \boldsymbol{\Sigma}_2) \\ & + p(\text{green})N(\mathbf{x} \mid \boldsymbol{\mu}_3, \boldsymbol{\Sigma}_3) \end{aligned}$$

# Learning GMMs

Learning a GMM means finding all the parameters  $\theta = \{\pi_j, \mu_j, \Sigma_j\}_{j=1}^k$ .

In the process, we will learn the distribution of the latent variable  $z_i$  as well:

$$p(z_i = j \mid \mathbf{x}_i) := \gamma_{ij} \in [0, 1]$$

i.e. “soft assignment” of each point to each cluster, as opposed to “hard assignment” by  $k$ -means.

GMM is more explanatory than  $k$ -means

- both learn the cluster centers  $\mu_j$ 's
- in addition, GMM learns cluster weight  $\pi_j$  and covariance  $\Sigma_j$ , thus
  - we can *predict probability of seeing a new point*
  - we can *generate synthetic data*

# Gaussian Mixture Model

- Introduction
- Learning the parameters
- EM algorithm
- EM for the Gaussian Mixture Model



# How do we learn the parameters?

As always, we want to do **maximum-likelihood estimation (MLE)**: find

$$\operatorname{argmax}_{\boldsymbol{\theta}} \ln \prod_{i=1}^n p(\mathbf{x}_i ; \boldsymbol{\theta}) = \operatorname{argmax}_{\boldsymbol{\theta}} \sum_{i=1}^n \ln p(\mathbf{x}_i ; \boldsymbol{\theta}) := \operatorname{argmax}_{\boldsymbol{\theta}} P(\boldsymbol{\theta}).$$

This is called **incomplete log-likelihood** (since  $z_i$ 's are unobserved). We can still write it down as an optimization problem by marginalizing out the  $z_i$ 's.

$$\begin{aligned} P(\boldsymbol{\theta}) &= \sum_{i=1}^n \ln p(\mathbf{x}_i ; \boldsymbol{\theta}) = \sum_{i=1}^n \ln \left( \sum_{j=1}^k p(\mathbf{x}_i, z_i = j ; \boldsymbol{\theta}) \right) \\ &= \sum_{i=1}^n \ln \left( \sum_{j=1}^k p(z_i = j ; \boldsymbol{\theta}) p(\mathbf{x}_i | z_i = j ; \boldsymbol{\theta}) \right) = \sum_{i=1}^n \ln \left( \sum_{j=1}^k \pi_j N(\mathbf{x}_i \mid \mu_j, \boldsymbol{\Sigma}_j) \right). \end{aligned}$$

This is a non-concave problem, and does not have a closed-form solution.

One solution is to still apply GD/SGD, but a much more effective approach is the **Expectation Maximization (EM) algorithm**.

# Preview of EM for learning GMMs

**Step 0** Initialize  $\pi_j, \mu_j, \Sigma_j$  for each  $j \in [k]$

**Step 1 (E-Step)** **update the “soft assignment”** (fixing parameters)

$$\gamma_{ij} = p(z_i = j \mid \mathbf{x}_i) \propto \pi_j N(\mathbf{x}_i \mid \mu_j, \Sigma_j)$$

**Step 2 (M-Step)** **update the model parameter** (fixing assignments)

$$\pi_j = \frac{\sum_i \gamma_{ij}}{n} \quad \mu_j = \frac{\sum_i \gamma_{ij} \mathbf{x}_i}{\sum_i \gamma_{ij}}$$

$$\Sigma_j = \frac{1}{\sum_i \gamma_{ij}} \sum_i \gamma_{ij} (\mathbf{x}_i - \mu_j)(\mathbf{x}_i - \mu_j)^T$$


**Step 3** return to Step 1 if not converged

We will see how this is **a special case of EM**.

# Demo

Iteration (t)

12



See Colab notebook

# Gaussian Mixture Model

- Introduction
- Learning the parameters
- **EM algorithm**
- EM for the Gaussian Mixture Model

# EM algorithm

In general EM is **a heuristic to solve MLE with latent variables** (not just GMM), i.e. find the maximizer of


$$P(\boldsymbol{\theta}) = \sum_{i=1}^n \ln p(\boldsymbol{x}_i ; \boldsymbol{\theta}) = \sum_{i=1}^n \ln \int_{z_i} p(\boldsymbol{x}_i, z_i ; \boldsymbol{\theta}) dz_i$$

- $\boldsymbol{\theta}$  is the **parameters** for a general probabilistic model
- $\boldsymbol{x}_i$ 's are **observed random variables**
- $z_i$ 's are **latent variables**

Again, directly solving the objective is usually complicated and does not have a closed form solution.

# High-level idea

Keep maximizing **a lower bound of  $P$**  that is more manageable



# Jensen's inequality

# A lower bound on the log likelihood

Finding the lower bound of  $P$ :

$$\begin{aligned}\ln p(\mathbf{x} ; \boldsymbol{\theta}) &= \ln \left( \sum_{z=1}^k p(\mathbf{x}, z ; \boldsymbol{\theta}) \right) \\ &= \ln \left( \sum_{z=1}^k q(z) \frac{p(\mathbf{x}, z ; \boldsymbol{\theta})}{q(z)} \right) \\ &= \ln \left( \mathbb{E}_{z \sim q(z)} \left[ \frac{p(\mathbf{x}, z ; \boldsymbol{\theta})}{q(z)} \right] \right) \\ &\geq \mathbb{E}_{z \sim q(z)} \left[ \ln \left( \frac{p(\mathbf{x}, z ; \boldsymbol{\theta})}{q(z)} \right) \right].\end{aligned}$$

Therefore, our log-likelihood can be written as,

$$\begin{aligned}P(\boldsymbol{\theta}) &= \sum_{i=1}^n \ln p(\mathbf{x}_i ; \boldsymbol{\theta}) \geq \sum_{i=1}^n \mathbb{E}_{z_i \sim q_i(z_i)} \left[ \ln \left( \frac{p(\mathbf{x}_i, z_i ; \boldsymbol{\theta})}{q_i(z_i)} \right) \right] \\ &= F(\boldsymbol{\theta}, \{q_i\}_{i=1}^n).\end{aligned}$$



## Alternatively maximizing the lower bound

The expression for the likelihood holds for *any*  $\{q_i\}$ , so how do we choose? If we have some guess of the parameters  $\theta$ , we should choose  $\{q_i\}$  to try to make the lower bound tight at that value of  $\theta$ , i.e. make the inequality hold with equality at that value of  $\theta$ .

Equivalently, this is the same as *alternatingly maximizing  $F$  over  $\{q_i\}$  and  $\theta$*  (similar to  $k$ -means).

Suppose we fix  $\theta^{(t)}$ , what should we choose  $\{q_i^{(t)}\}$ ?

The inequality arises from the step where we used Jensen's inequality. How do we get this step to hold with equality? The function should be a constant function, i.e.

$$\frac{p(\mathbf{x}_i, z_i ; \theta^{(t)})}{q_i^{(t)}(z_i)} = c_i$$

for some constant  $c_i$  which does not depend on the value taken by the random variable  $z_i$ .

## Maximizing over $\{q_i\}$

(continued) Since  $\sum_{z_i=1}^k q_i^{(t)}(z_i) = 1$ , we get,

$$\begin{aligned} c_i &= \sum_{z_i=1}^k p(\mathbf{x}_i, z_i ; \boldsymbol{\theta}^{(t)}) \\ \implies q_i^{(t)}(z_i) &= \frac{p(\mathbf{x}_i, z_i ; \boldsymbol{\theta}^{(t)})}{\sum_{z_i=1}^k p(\mathbf{x}_i, z_i ; \boldsymbol{\theta}^{(t)})} \\ &= \frac{p(\mathbf{x}_i, z_i ; \boldsymbol{\theta}^{(t)})}{p(\mathbf{x}_i ; \boldsymbol{\theta}^{(t)})} \\ &= p(z_i | \mathbf{x}_i ; \boldsymbol{\theta}^{(t)}) \end{aligned}$$

i.e., the *posterior distribution of  $z_i$*  given  $\mathbf{x}_i$  and  $\boldsymbol{\theta}^{(t)}$ .

So at  $\boldsymbol{\theta}^{(t)}$ , we found the tightest lower bound  $F(\boldsymbol{\theta}, \{q_i^{(t)}\})$ :

- $F(\boldsymbol{\theta}, \{q_i^{(t)}\}) \leq P(\boldsymbol{\theta})$  for all  $\boldsymbol{\theta}$ .
- $F(\boldsymbol{\theta}^{(t)}, \{q_i^{(t)}\}) = P(\boldsymbol{\theta}^{(t)})$

# Maximizing over $\theta$

Fix  $\{q_i^{(t)}\}$ , maximize over  $\theta$ :

$$\begin{aligned} & \operatorname{argmax}_{\theta} F(\theta, \{q_i^{(t)}\}) \\ &= \operatorname{argmax}_{\theta} \sum_{i=1}^n \mathbb{E}_{z_i \sim q_i^{(t)}} \left[ \ln \left( \frac{p(\mathbf{x}_i, z_i; \theta)}{q_i^{(t)}(z_i)} \right) \right] \\ &= \operatorname{argmax}_{\theta} \sum_{i=1}^n \mathbb{E}_{z_i \sim q_i^{(t)}} [\ln p(\mathbf{x}_i, z_i; \theta)] - \sum_{i=1}^n \mathbb{E}_{z_i \sim q_i^{(t)}} [\ln(q_i^{(t)}(z_i))] \\ &= \operatorname{argmax}_{\theta} \sum_{i=1}^n \mathbb{E}_{z_i \sim q_i^{(t)}} [\ln p(\mathbf{x}_i, z_i; \theta)] \\ &:= \operatorname{argmax}_{\theta} Q(\theta; \theta^{(t)}) \end{aligned} \quad (\{q_i^{(t)}\} \text{ are computed via } \theta^{(t)})$$

$Q$  is the (expected) **complete likelihood** and is usually more tractable.

- versus the incomplete likelihood:  $P(\theta) = \sum_{i=1}^n \ln p(\mathbf{x}_i; \theta)$

# General EM algorithm

**Step 0** Initialize  $\theta^{(1)}$ ,  $t = 1$

**Step 1 (E-Step)** **update the posterior of latent variables**  $z_i$ ,

$$q_i^{(t)}(z_i) = p(z_i \mid \mathbf{x}_i ; \theta^{(t)})$$

and obtain **Expectation** of complete likelihood


$$Q(\theta ; \theta^{(t)}) = \sum_{i=1}^n \mathbb{E}_{z_i \sim q_i^{(t)}} [\ln p(\mathbf{x}_i, z_i ; \theta)]$$

**Step 2 (M-Step)** **update the model parameter** via **Maximization**

$$\theta^{(t+1)} \leftarrow \underset{\theta}{\operatorname{argmax}} Q(\theta ; \theta^{(t)})$$

**Step 3**  $t \leftarrow t + 1$  and return to Step 1 if not converged

# Pictorial explanation



$P(\theta)$  is non-concave, but  $Q(\theta; \theta^{(t)})$  often is concave and easy to maximize.

$$\begin{aligned} P(\theta^{(t+1)}) &\geq F(\theta^{(t+1)}; \{q_i^{(t)}\}) \\ &\geq F(\theta^{(t)}; \{q_i^{(t)}\}) \\ &= P(\theta^{(t)}) \end{aligned}$$

So **EM** always **increases the objective value** and will **converge to some local maximum** (similar to  $k$ -means).

# Gaussian Mixture Model

- Introduction
- Learning the parameters
- EM algorithm
- EM for the Gaussian Mixture Model

# Applying EM to learn GMMs: E-Step

**E-Step:**

$$\begin{aligned} q_i^{(t)}(z_i = j) &= p(z_i = j \mid \mathbf{x}_i ; \boldsymbol{\theta}^{(t)}) \\ &\propto p(\mathbf{x}_i, z_i = j ; \boldsymbol{\theta}^{(t)}) \\ &= p(z_i = j ; \boldsymbol{\theta}^{(t)}) p(\mathbf{x}_i \mid z_i = j ; \boldsymbol{\theta}^{(t)}) \\ &= \pi_j^{(t)} N(\mathbf{x}_i \mid \boldsymbol{\mu}_j^{(t)}, \boldsymbol{\Sigma}_j^{(t)}) \end{aligned}$$

This computes the “soft assignment”  $\gamma_{ij} = q_i^{(t)}(z_i = j)$ , i.e. conditional probability of  $\mathbf{x}_i$  belonging to cluster  $k$ .

# Applying EM to learn GMMs: M-Step

**M-Step:**

$$\begin{aligned}\operatorname{argmax}_{\boldsymbol{\theta}} Q(\boldsymbol{\theta}, \boldsymbol{\theta}^{(t)}) &= \operatorname{argmax}_{\boldsymbol{\theta}} \sum_{i=1}^n \mathbb{E}_{z_i \sim q_i^{(t)}} [\ln p(\mathbf{x}_i, z_i ; \boldsymbol{\theta})] \\ &= \operatorname{argmax}_{\boldsymbol{\theta}} \sum_{i=1}^n \mathbb{E}_{z_i \sim q_i^{(t)}} [\ln p(z_i ; \boldsymbol{\theta}) + \ln p(\mathbf{x}_i | z_i ; \boldsymbol{\theta})] \\ &= \operatorname{argmax}_{\{\pi_j, \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j\}} \sum_{i=1}^n \sum_{j=1}^k \gamma_{ij} (\ln \pi_j + \ln N(\mathbf{x}_i | \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j))\end{aligned}$$

To find  $\pi_1, \dots, \pi_k$ , solve

$$\operatorname{argmax}_{\boldsymbol{\pi}} \sum_{i=1}^n \sum_{j=1}^k \gamma_{ij} \ln \pi_j$$

To find each  $\boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j$ , solve

$$\operatorname{argmax}_{\boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j} \sum_{i=1}^n \gamma_{ij} \ln N(\mathbf{x}_i | \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)$$



# Applying EM to learn GMMs: M-Step

Solutions to previous two problems are very natural, for each  $j$

$$\pi_j = \frac{\sum_i \gamma_{ij}}{n}$$

i.e. (weighted) fraction of examples belonging to cluster  $j$

$$\boldsymbol{\mu}_j = \frac{\sum_i \gamma_{ij} \mathbf{x}_i}{\sum_i \gamma_{ij}}$$

i.e. (weighted) average of examples belonging to cluster  $j$

$$\boldsymbol{\Sigma}_j = \frac{1}{\sum_i \gamma_{ij}} \sum_i \gamma_{ij} (\mathbf{x}_i - \boldsymbol{\mu}_j)(\mathbf{x}_i - \boldsymbol{\mu}_j)^T$$

i.e (weighted) covariance of examples belonging to cluster  $j$

You will verify some of these in HW4.

# Applying EM to learn GMMs: Putting it together

EM for learning GMMs:

**Step 0** Initialize  $\pi_j, \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j$  for each  $j \in [k]$

**Step 1 (E-Step)** **update the “soft assignment”** (fixing parameters)

$$\gamma_{ij} = p(z_i = j \mid \mathbf{x}_i) \propto \pi_j N(\mathbf{x}_i \mid \boldsymbol{\mu}_j, \boldsymbol{\Sigma}_j)$$

**Step 2 (M-Step)** **update the model parameter** (fixing assignments)

$$\pi_j = \frac{\sum_i \gamma_{ij}}{n} \quad \boldsymbol{\mu}_j = \frac{\sum_i \gamma_{ij} \mathbf{x}_i}{\sum_i \gamma_{ij}}$$

$$\boldsymbol{\Sigma}_j = \frac{1}{\sum_i \gamma_{ij}} \sum_i \gamma_{ij} (\mathbf{x}_i - \boldsymbol{\mu}_j)(\mathbf{x}_i - \boldsymbol{\mu}_j)^T$$

**Step 3** return to Step 1 if not converged